

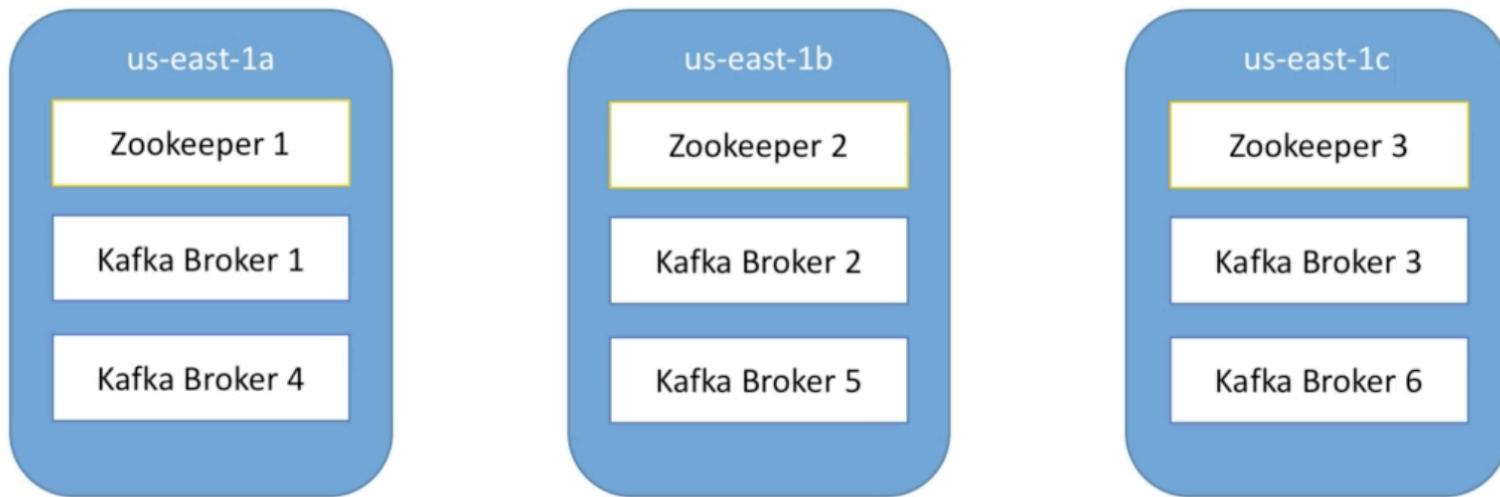
Kafka Cluster Setup

High Level Architecture



© Stephane Maarek

- You want multiple brokers in different data centers (racks) to distribute your load. You also want a cluster of at least 3 zookeeper
- In AWS:





© Stephane Maarek

Kafka Cluster Setup Gotchas

- It's not easy to setup a cluster
- You want to isolate each Zookeeper & Broker on separate servers
- Monitoring needs to be implemented
- Operations have to be mastered
- You need a really good Kafka Admin

- Alternative: many different “Kafka as a Service” offerings on the web
- No operational burdens (updates, monitoring, setup, etc...)



Kafka Monitoring and Operations

- Kafka exposes metrics through JMX.
- These metrics are highly important for monitoring Kafka, and ensuring the systems are behaving correctly under load.
- Common places to host the Kafka metrics:
 - ELK (ElasticSearch + Kibana)
 - Datadog
 - NewRelic
 - Confluent Control Centre
 - Prometheus
 - Many others...!

Kafka Monitoring and Operations



© Stephane Maarek

- Some of the most important metrics are:
- Under Replicated Partitions: Number of partitions are have problems with the ISR (in-sync replicas). May indicate a high load on the system
- Request Handlers: utilization of threads for IO, network, etc... overall utilization of an Apache Kafka broker.
- Request timing: how long it takes to reply to requests. Lower is better, as latency will be improved.



Kafka Monitoring and Operations

- Overall have a look at the documentation here:
- <https://kafka.apache.org/documentation/#monitoring>
- <https://docs.confluent.io/current/kafka/monitoring.html>
- <https://www.datadoghq.com/blog/monitoring-kafka-performance-metrics/>



Kafka Monitoring and Operations

- Kafka Operations team must be able to perform the following tasks:
 - Rolling Restart of Brokers
 - Updating Configurations
 - Rebalancing Partitions
 - Increasing replication factor
 - Adding a Broker
 - Replacing a Broker
 - Removing a Broker
 - Upgrading a Kafka Cluster with zero downtime

The need for encryption, authentication & authorization in Kafka



- Currently, any client can access your Kafka cluster (**authentication**)
- The clients can publish / consume any topic data (**authorisation**)
- All the data being sent is fully visible on the network (**encryption**)

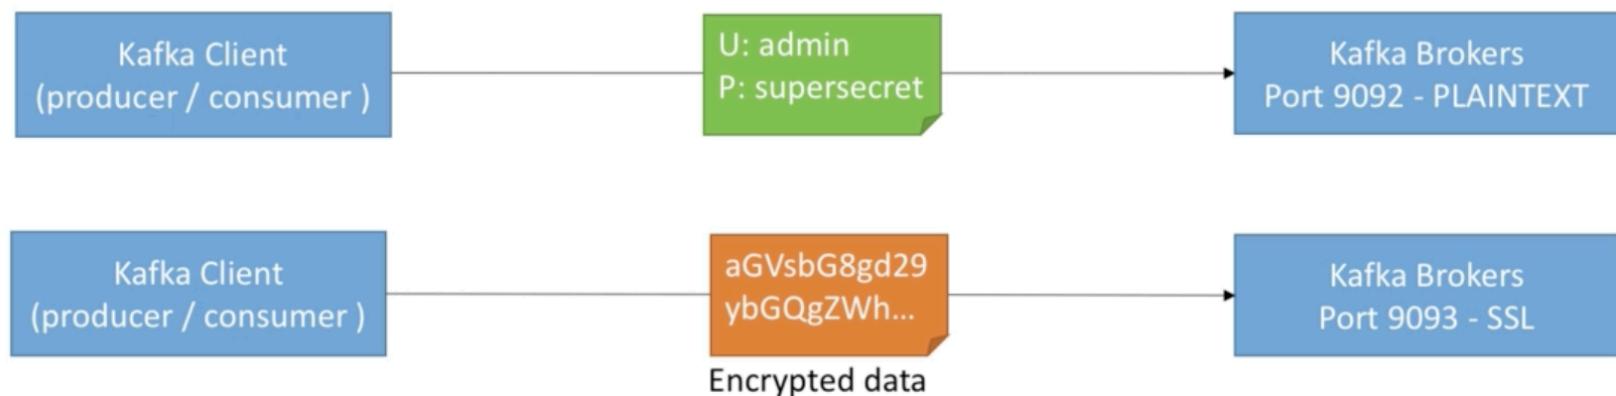
- Someone could **intercept data being sent**
- Someone could **publish bad data / steal data**
- Someone could **delete topics**

- All these reasons push for more security and an authentication model



Encryption in Kafka

- Encryption in Kafka ensures that the data exchanged between clients and brokers is secret to routers on the way
- This is similar concept to an https website

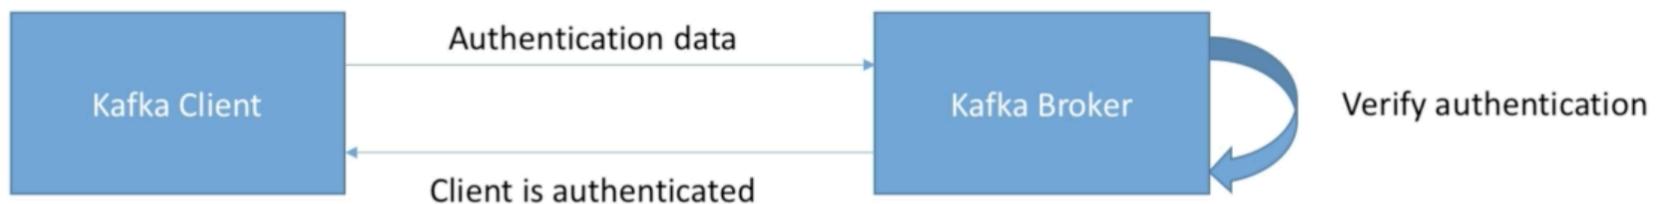




© Stephane Maarek

Authentication in Kafka

- Authentication in Kafka ensures that only **clients that can prove their identity** can connect to our Kafka Cluster
- This is similar concept to a login (username / password)





Authentication in Kafka

- Authentication in Kafka can take a few forms
- SSL Authentication: clients authenticate to Kafka using SSL certificates
- SASL Authentication:
 - PLAIN: clients authenticate using username / password (weak – easy to setup)
 - Kerberos: such as Microsoft Active Directory (strong – hard to setup)
 - SCRAM: username / password (strong – medium to setup)



Authorisation in Kafka

- Once a client is authenticated, Kafka can verify its identity
- It still needs to be combined with authorisation, so that Kafka knows that
 - "User alice can view topic finance"
 - "User bob cannot view topic trucks"
- ACL (Access Control Lists) have to be maintained by administration and onboard new users

Putting it all together



- You can mix
 - Encryption
 - Authentication
 - Authorisation
- This allows your Kafka clients to:
 - Communicate securely to Kafka
 - Clients would authenticate against Kafka
 - Kafka can authorise clients to read / write to topics

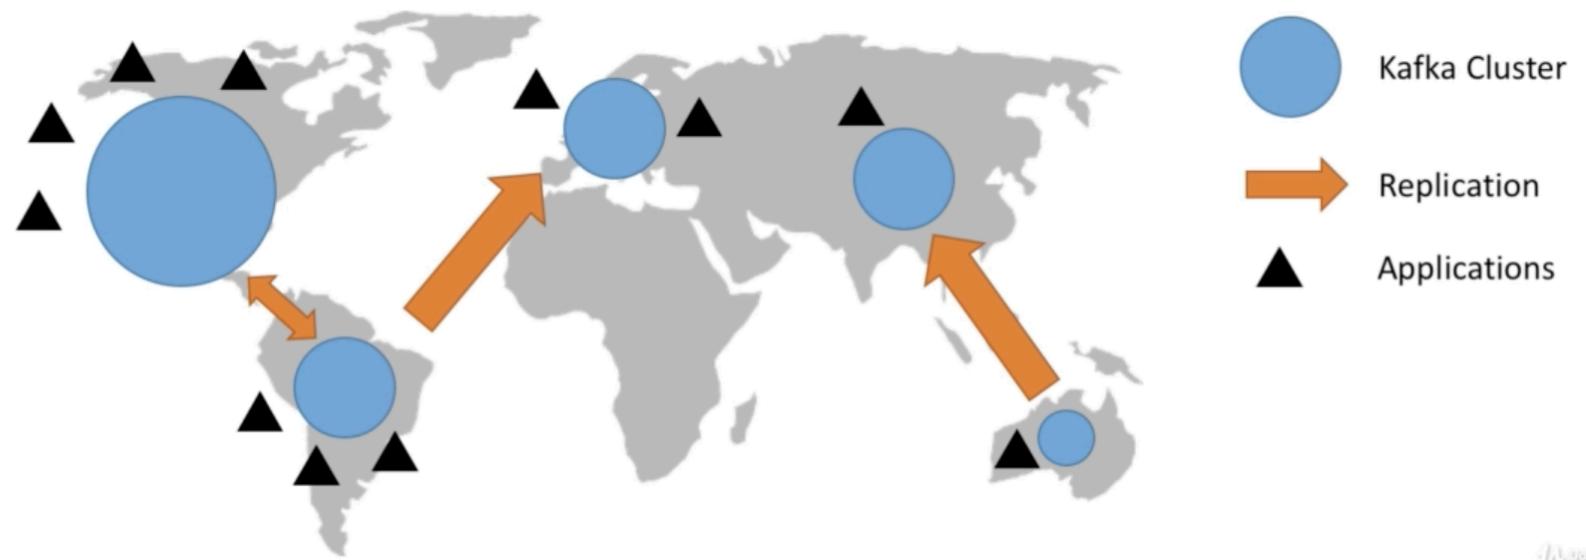
State of the art of Kafka Security



- Kafka Security is fairly new (0.10)
- Kafka Security improves over time and becomes more flexible / easier to setup as time goes.
- Currently, it is **hard** to setup Kafka Security.
- Best support for Kafka Security for applications is with **Java**.

Kafka Multi Cluster + Replication

- Kafka can only operate well in a single region
- Therefore, it is very common for enterprises to have Kafka clusters across the world, with some level of replication between them



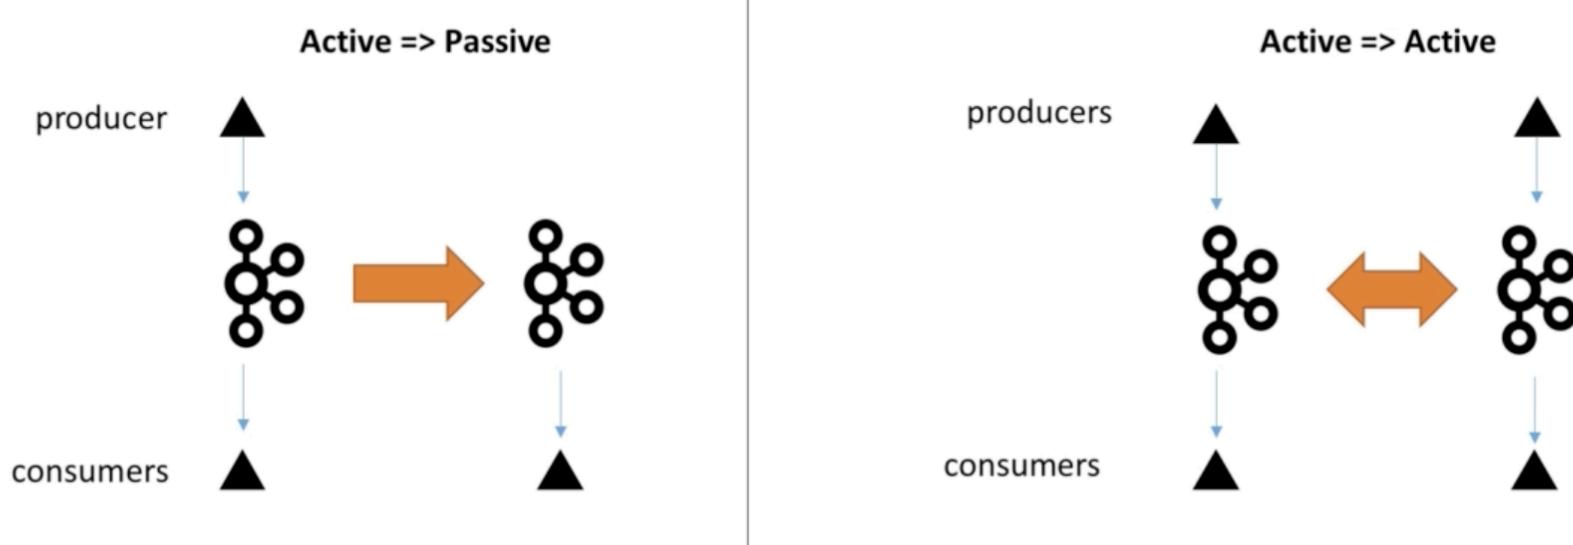


Kafka Multi Cluster + Replication

- A replication application at its core is just a consumer + a producer
- There are different tools to perform it:
 - Mirror Maker - open source tool that ships with Kafka
 - Netflix uses Flink – they wrote their own application
 - Uber uses uReplicator – addresses performance and operations issues with MM
 - Comcast has their own open source Kafka Connect Source
 - Confluent has their own Kafka Connect Source (paid)
- Overall, try these and see if it works for your use case before writing your own

Kafka Multi Cluster + Replication

- There are two designs for cluster replication:





Kafka Multi Cluster + Replication

- Active => Active:
 - You have a global application
 - You have a global dataset
- Active => Passive:
 - You want to have an aggregation cluster (for example for analytics)
 - You want to create some form of disaster recovery strategy (*it's hard)
 - Cloud Migration (from on-premise cluster to Cloud cluster)
- **Replicating doesn't preserve offsets, just data!**

Kafka Multi Cluster + Replication



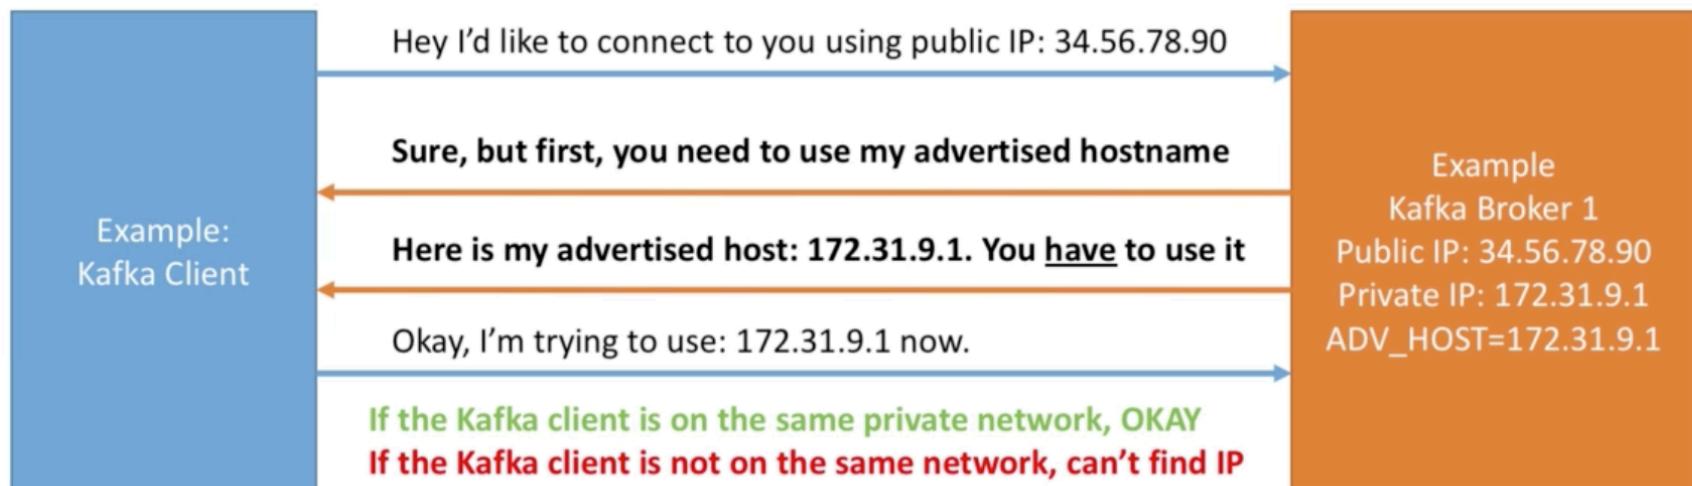
© Stephane Maarek

- Optimization resources online:
 - <https://cwiki.apache.org/confluence/pages/viewpage.action?pageId=27846330>
 - <https://engineering.salesforce.com/mirrormaker-performance-tuning-63afaed12c21>
 - <https://docs.confluent.io/current/multi-dc/replicator-tuning.html#improving-network-utilization-of-a-connect-task>
 - <https://community.hortonworks.com/articles/79891/kafka-mirror-maker-best-practices.html>
 - <https://www.confluent.io/kafka-summit-sf17/multitenant-multicloud-and-hierarchical-kafka-messaging-service>
 - <https://eng.uber.com/ureplicator/>
 - <https://www.altoros.com/blog/multi-cluster-deployment-options-for-apache-kafka-pros-and-cons/>
 - <https://github.com/Comcast/MirrorTool-for-Kafka-Connect>



Understanding communications between Client and with Kafka

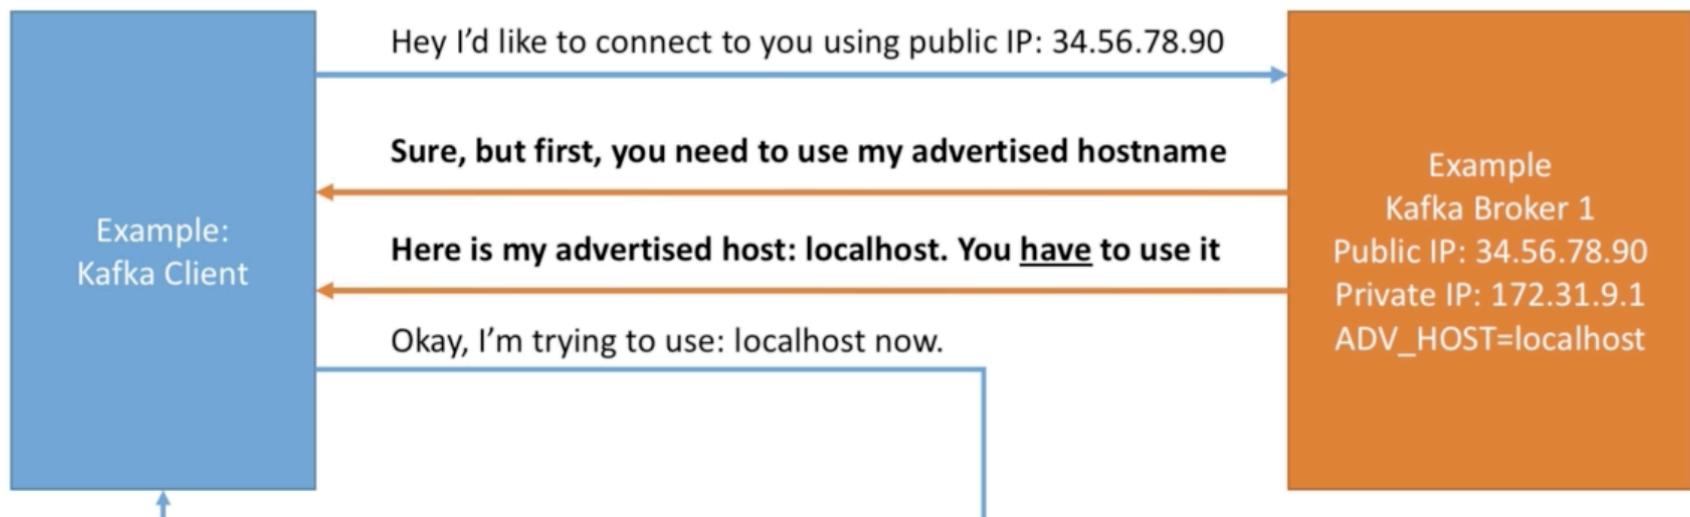
- **Advertised listeners is the most important setting of Kafka**





But what if I put localhost? It works on my machine!

- **Advertised hostname is the most important setting of Kafka**



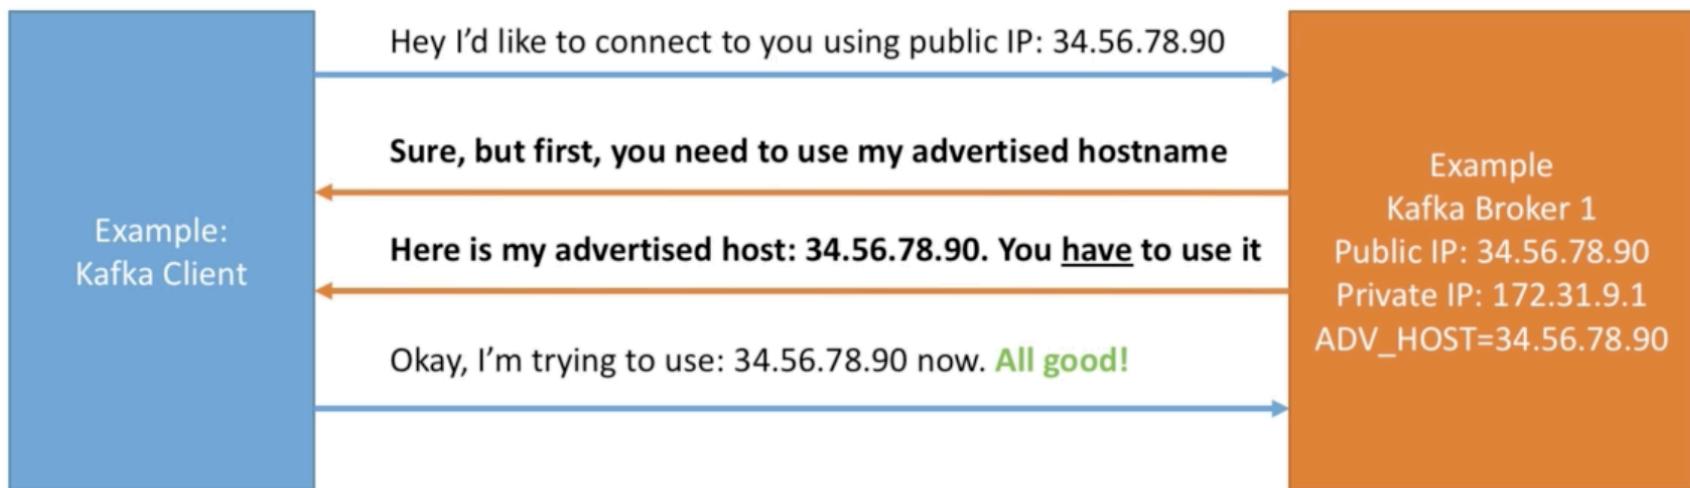
If the Kafka client is on the machine and you have 1 broker, OKAY

If the Kafka client is not on the same machine or you have 2 or more broker, NOT OKAY



What if I use the public IP?

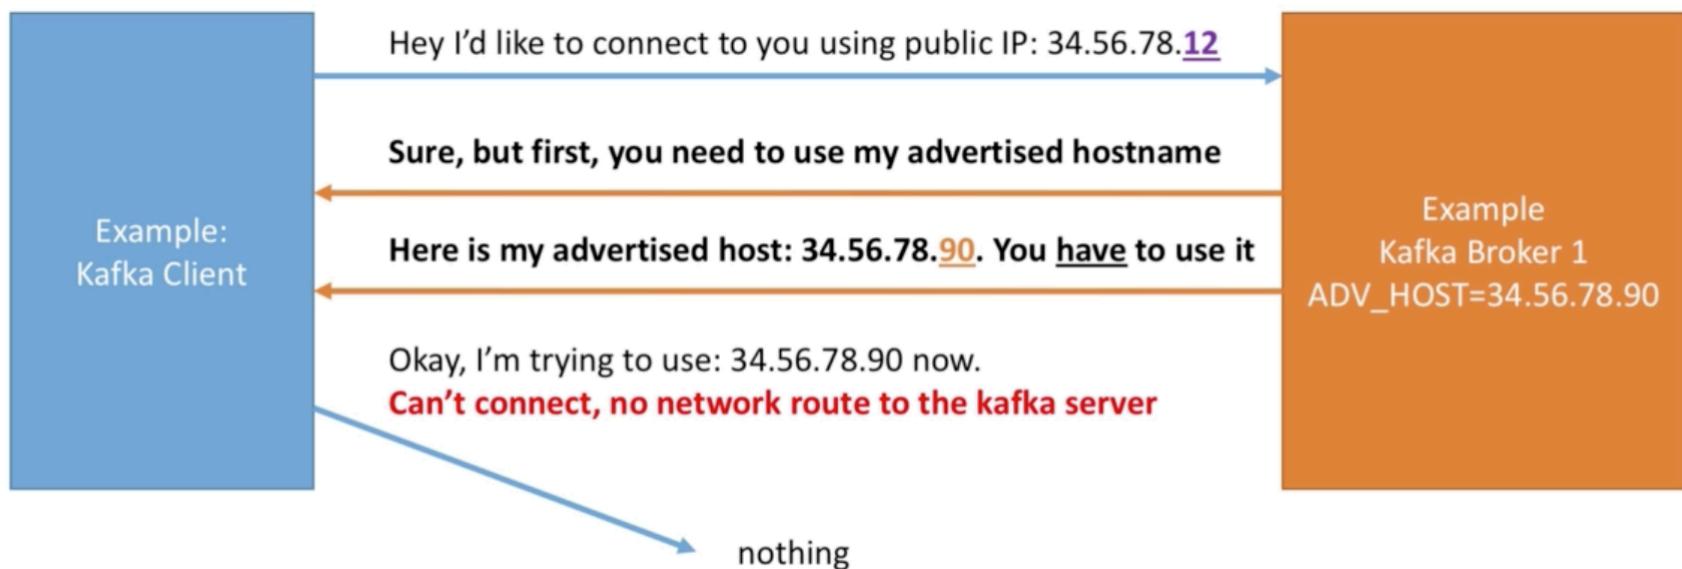
- **Advertised hostname is the most important setting of Kafka**





What if I use the public IP... But after a reboot Kafka public IP changed!

- Assume the IP of your server has changed:
 - FROM 34.56.78.90 => TO 34.56.78.12



So what do I set for advertised.listeners?



- If your clients are on your private network, set either:
 - the internal private IP
 - the internal private DNS hostname
- Your clients should be able to resolve the internal IP or hostname
- If your clients are on a public network, set either:
 - The external public IP
 - The external public DNS hostname pointing to the public IP
- Your clients must be able to resolve the public DNS